

Javascript Bignum Extensions

Version 2020-01-11

Author: Fabrice Bellard

Table of Contents

1	Introduction	1
2	Operator overloading	2
3	BigInt extensions	3
4	BigFloat	4
4.1	Introduction	4
4.2	Floating point rounding	4
4.3	Operators	4
4.4	BigFloat literals	4
4.5	Builtin Object changes	5
4.5.1	BigFloat function	5
4.5.2	BigFloat.prototype	6
4.5.3	BigFloatEnv constructor	6
5	BigDecimal	9
5.1	Operators	9
5.2	BigDecimal literals	9
5.3	Builtin Object changes	9
5.3.1	The BigDecimal function	9
5.3.2	Properties of the BigDecimal object	10
5.3.3	Properties of the BigDecimal.prototype object	10
6	Math mode	11

1 Introduction

The Bignum extensions add the following features to the Javascript language while being 100% backward compatible:

- Operator overloading with a dispatch logic inspired from the proposal available at <https://github.com/tc39/proposal-operator-overloading/>.

- Arbitrarily large floating point numbers (`BigFloat`) in base 2 using the IEEE 754 semantics.

- Arbitrarily large floating point numbers (`BigDecimal`) in base 10 based on the proposal available at <https://github.com/littledan/proposal-bigdecimal>.

- math mode: arbitrarily large integers and floating point numbers are available by default. The integer division and power can be overloaded for example to return a fraction. The modulo operator (%) is defined as the Euclidian remainder. ^ is an alias to the power operator (**). ^^ is used as the exclusive or operator.

The extensions are independent from each other except the math mode which relies on `BigFloat` and operator overloading.

2 Operator overloading

Operator overloading is inspired from the proposal available at <https://github.com/tc39/proposal-operator-overloading/>. It implements the same dispatch logic but finds the operator sets by looking at the `Symbol.operatorSet` property in the objects. The changes were done in order to simplify the implementation.

More precisely, the following modifications were made:

- `with` operators from is not supported. Operator overloading is always enabled.

- The dispatch is not based on a static `[[OperatorSet]]` field in all instances. Instead, a dynamic lookup of the `Symbol.operatorSet` property is done. This property is typically added in the prototype of each object.

- `Operators.create(...dictionaries)` is used to create a new `OperatorSet` object. The `Operators` function is supported as an helper to be closer to the TC39 proposal.

- `[]` cannot be overloaded.

- In math mode, the `BigInt` division and power operators can be overloaded with `Operators.updateBigIntOperators(dictionary)`.

3 BigInt extensions

A few properties are added to the BigInt object:

- tdi v(a, b) Return $\text{trunc}(a/b)$. $b = 0$ raises a RangeError exception.
- fdi v(a, b) Return ba/bc . $b = 0$ raises a RangeError exception.
- cdi v(a, b) Return da/be . $b = 0$ raises a RangeError exception.
- edi v(a, b) Return $\text{sgn}(b)ba/\text{bjc}$ (Euclidian division). $b = 0$ raises a RangeError exception.
- tdi vrem(a, b)
fdi vrem(a, b)
cdi vrem(a, b)
edi vrem(a, b) Return an array of two elements. The first element is the quotient, the second is the remainder. The same rounding is done as the corresponding division operation.
- sqrt(a) Return $b^{\frac{1}{2}}(a)c$. A RangeError exception is raised if $a < 0$.
- sqrtrem(a) Return an array of two elements. The first element is $b^{\frac{1}{2}}ac$. The second element is $a - b^{\frac{1}{2}}ac^2$. A RangeError exception is raised if $a < 0$.
- floorLog2(a) Return -1 if $a = 0$ otherwise return $\text{blog}_2(a)c$.
- ctz(a) Return the number of trailing zeros in the two's complement binary representation of a . Return -1 if $a = 0$.

4 BigFloat

4.1 Introduction

This extension adds the `BigFloat` primitive type. The `BigFloat` type represents floating point numbers in base 2 with the IEEE 754 semantics. A floating point number is represented as a sign, mantissa and exponent. The special values NaN, $+\infty$, $-\infty$, +0 and -0 are supported. The mantissa and exponent can have any bit length with an implementation specific minimum and maximum.

4.2 Floating point rounding

Each floating point operation operates with infinite precision and then rounds the result according to the specified floating point environment (`BigFloatEnv` object). The status flags of the environment are also set according to the result of the operation.

If no floating point environment is provided, the global floating point environment is used.

The rounding mode of the global floating point environment is always RNDN ("round to nearest with ties to even")¹. The status flags of the global environment cannot be read². The precision of the global environment is `BigFloatEnv.prec`. The number of exponent bits of the global environment is `BigFloatEnv.expBits`. The global environment subnormal flag is set to true.

For example, `prec = 53` and `expBits = 11` exactly give the same precision as the IEEE 754 64 bit floating point format. The default precision is `prec = 113` and `expBits = 15` (IEEE 754 128 bit floating point format).

The global floating point environment can only be modified temporarily when calling a function (see `BigFloatEnv.setPrec`). Hence a function can change the global floating point environment for its callees but not for its caller.

4.3 Operators

The builtin operators are extended so that a `BigFloat` is returned if at least one operand is a `BigFloat`. The computations are always done with infinite precision and rounded according to the global floating point environment.

`typeof` applied on a `BigFloat` returns `bigfloat`.

`BigFloat` can be compared with all the other numeric types and the result follows the expected mathematical relations.

However, since `BigFloat` and `Number` are different types they are never equal when using the strict comparison operators (e.g. `0.0 === 0.01` is false).

4.4 BigFloat literals

`BigFloat` literals are floating point numbers with a trailing `L` suffix. `BigFloat` literals have infinite precision. They are rounded according to the global floating point environment when they are evaluated.³

¹ The rationale is that the rounding mode changes must always be explicit.

² The rationale is to avoid side effects for the built-in operators.

³ Base 10 floating point literals cannot usually be exactly represented as base 2 floating point number. In order to ensure that the literal is represented accurately with the current precision, it must be evaluated at runtime.

4.5 Builtin Object changes

4.5.1 BigFloat function

The BigFloat function cannot be invoked as a constructor. When invoked as a function: the parameter is converted to a primitive type. If the result is a numeric type, it is converted to BigFloat without rounding. If the result is a string, it is converted to BigFloat using the precision of the global floating point environment.

BigFloat properties:

LN2

PI Getter. Return the value of the corresponding mathematical constant rounded to nearest, ties to even with the current global precision. The constant values are cached for small precisions.

MIN_VALUE

MAX_VALUE

EPSILON Getter. Return the minimum, maximum and epsilon BigFloat values (same definition as the corresponding Number constants).

fpRound(a[, e])

Round the floating point number a according to the floating point environment e or the global environment if e is undefined.

parseFloat(a[, radix[, e]])

Parse the string a as a floating point number in radix radix. The radix is 0 (default) or from 2 to 36. The radix 0 means radix 10 unless there is a hexadecimal or binary prefix. The result is rounded according to the floating point environment e or the global environment if e is undefined.

isFinite(a)

Return true if a is a finite big float.

isNaN(a) Return true if a is a NaN big float.

add(a, b[, e])

sub(a, b[, e])

mul(a, b[, e])

div(a, b[, e])

Perform the specified floating point operation and round the floating point number a according to the floating point environment e or the global environment if e is undefined. If e is specified, the floating point status flags are updated.

floor(x)

ceil(x)

round(x)

trunc(x) Round to an integer. No additional rounding is performed.

abs(x) Return the absolute value of x. No additional rounding is performed.

fmod(x, y[, e])

remainder(x, y[, e])

Floating point remainder. The quotient is truncated to zero (fmod) or to the nearest integer with ties to even (remainder). e is an optional floating point environment.

sqrt(x[, e])

Square root. Return a rounded floating point number. e is an optional floating point environment.

```

sin(x[, e])
cos(x[, e])
tan(x[, e])
asin(x[, e])
acos(x[, e])
atan(x[, e])
atan2(x, y[, e])
exp(x[, e])
log(x[, e])
pow(x, y[, e])

```

Transcendental operations. Return a rounded floating point number. *e* is an optional floating point environment.

4.5.2 BigFloat.prototype

The following properties are modified:

`valueOf()`

Return the big float primitive value corresponding to this.

`toString(radix)`

For floating point numbers:

If the radix is a power of two, the conversion is done with infinite precision.

Otherwise, the number is rounded to nearest with ties to even using the global precision. It is then converted to string using the minimum number of digits so that its conversion back to a floating point using the global precision and round to nearest gives the same number.

The exponent letter is *e* for base 10, *p* for bases 2, 8, 16 with a binary exponent and *@* for the other bases.

`toPrecision(p, rnd_mode = BigFloatEnv.RNDNA, radix = 10)`

`toFixed(p, rnd_mode = BigFloatEnv.RNDNA, radix = 10)`

`toExponential(p, rnd_mode = BigFloatEnv.RNDNA, radix = 10)`

Same semantics as the corresponding Number functions with BigFloats. There is no limit on the accepted precision *p*. The rounding mode and radix can be optionally specified. The radix must be between 2 and 36.

4.5.3 BigFloatEnv constructor

The `BigFloatEnv([p, [, rndMode]])` constructor cannot be invoked as a function. The floating point environment contains:

the mantissa precision in bits

the exponent size in bits assuming an IEEE 754 representation;

the subnormal flag (if true, subnormal floating point numbers can be generated by the floating point operations).

the rounding mode

the floating point status. The status flags can only be set by the floating point operations. They can be reset with `BigFloatEnv.prototype.clearStatus()` or with the various status flag setters.

`new BigFloatEnv([p, [, rndMode]])` creates a new floating point environment. The status flags are reset. If no parameter is given the precision, exponent bits and subnormal flags are copied from the global floating point environment. Otherwise, the precision is set to *p*, the

number of exponent bits is set to `expBitsMax` and the `subnormal` flag is set to `false`. If `rndMode` is undefined, the rounding mode is set to `RNDN`.

`BigFloatEnv` properties:

<code>prec</code>	Getter. Return the precision in bits of the global floating point environment. The initial value is 113.
<code>expBits</code>	Getter. Return the exponent size in bits of the global floating point environment assuming an IEEE 754 representation. The initial value is 15.
<code>setPrec(f, p[, e])</code>	Set the precision of the global floating point environment to <code>p</code> and the exponent size to <code>e</code> then call the function <code>f</code> . Then the <code>Float</code> precision and exponent size are reset to their previous value and the return value of <code>f</code> is returned (or an exception is raised if <code>f</code> raised an exception). If <code>e</code> is undefined it is set to <code>BigFloatEnv.expBitsMax</code> .
<code>precMin</code>	Read-only integer. Return the minimum allowed precision. Must be at least 2.
<code>precMax</code>	Read-only integer. Return the maximum allowed precision. Must be at least 113.
<code>expBitsMin</code>	Read-only integer. Return the minimum allowed exponent size in bits. Must be at least 3.
<code>expBitsMax</code>	Read-only integer. Return the maximum allowed exponent size in bits. Must be at least 15.
<code>RNDN</code>	Read-only integer. Round to nearest, with ties to even rounding mode.
<code>RNDZ</code>	Read-only integer. Round to zero rounding mode.
<code>RNDD</code>	Read-only integer. Round to -Infinity rounding mode.
<code>RNDU</code>	Read-only integer. Round to +Infinity rounding mode.
<code>RNDNA</code>	Read-only integer. Round to nearest, with ties away from zero rounding mode.
<code>RNDA</code>	Read-only integer. Round away from zero rounding mode.
<code>RNDF</code> ⁴	Read-only integer. Faithful rounding mode. The result is non-deterministically rounded to -Infinity or +Infinity. This rounding mode usually gives a faster and deterministic running time for the floating point operations.

`BigFloatEnv` prototype properties:

<code>prec</code>	Getter and setter (Integer). Return or set the precision in bits.
<code>expBits</code>	Getter and setter (Integer). Return or set the exponent size in bits assuming an IEEE 754 representation.
<code>rndMode</code>	Getter and setter (Integer). Return or set the rounding mode.
<code>subnormal</code>	Getter and setter (Boolean). <code>subnormal</code> flag. It is false when <code>expBits = expBitsMax</code> .
<code>clearStatus()</code>	Clear the status flags.
<code>invalidOperation</code>	
<code>divideByZero</code>	
<code>overflow</code>	

⁴ Could be removed in case a deterministic behavior for floating point operations is required.

underflow

inexact Getter and setter (Boolean). Status flags.

5 BigDecimal

This extension adds the `BigDecimal` primitive type. The `BigDecimal` type represents floating point numbers in base 10. It is inspired from the proposal available at <https://github.com/littledan/proposal-bigdecimal>.

The `BigDecimal` floating point numbers are always normalized and finite. There is no concept of `-0`, `Infinity` or `NaN`. By default, all the computations are done with infinite precision.

5.1 Operators

The following builtin operators support `BigDecimal`:

<code>+</code>	
<code>-</code>	
<code>*</code>	Both operands must be <code>BigDecimal</code> . The result is computed with infinite precision.
<code>%</code>	Both operands must be <code>BigDecimal</code> . The result is computed with infinite precision. A range error is thrown in case of division by zero.
<code>/</code>	Both operands must be <code>BigDecimal</code> . A range error is thrown in case of division by zero or if the result cannot be represented with infinite precision (use <code>BigDecimal.div</code> to specify the rounding).
<code>**</code>	Both operands must be <code>BigDecimal</code> . The exponent must be a positive integer. The result is computed with infinite precision.
<code>===</code>	When one of the operand is a <code>BigDecimal</code> , return true if both operands are a <code>BigDecimal</code> and if they are equal.
<code>==</code>	
<code>!=</code>	
<code><=</code>	
<code>>=</code>	
<code><</code>	
<code>></code>	Numerical comparison. When one of the operand is not a <code>BigDecimal</code> , it is converted to <code>BigDecimal</code> by using <code>ToString()</code> . Hence comparisons between <code>Number</code> and <code>BigDecimal</code> do not use the exact mathematical value of the <code>Number</code> value.

5.2 BigDecimal literals

`BigDecimal` literals are decimal floating point numbers with a trailing `m` suffix.

5.3 Builtin Object changes

5.3.1 The `BigDecimal` function.

It returns `0m` if no parameter is provided. Otherwise the first parameter is converted to a `BigDecimal` by using `ToString()`. Hence `Number` values are not converted to their exact numerical value as `BigDecimal`.

5.3.2 Properties of the BigDecimal object

`add(a, b[, e])`
`sub(a, b[, e])`
`mul(a, b[, e])`
`div(a, b[, e])`
`mod(a, b[, e])`
`sqrt(a, e)`
`round(a, e)`

Perform the specified floating point operation and round the floating point result according to the rounding object `e`. If the rounding object is not present, the operation is executed with infinite precision.

For `div`, a `RangeError` exception is thrown in case of division by zero or if the result cannot be represented with infinite precision if no rounding object is present.

For `sqrt`, a range error is thrown if `a` is less than zero.

The rounding object must contain the following properties: `roundingMode` is a string specifying the rounding mode ("floor", "ceiling", "down", "up", "half-even", "half-up"). Either `maximumSignificantDigits` or `maximumFractionDigits` must be present to specify respectively the number of significant digits (must be ≥ 1) or the number of digits after the decimal point (must be ≥ 0).

5.3.3 Properties of the BigDecimal.prototype object

`valueOf()`

Return the bigdecimal primitive value corresponding to this.

`toString()`

Convert this to a string with infinite precision in base 10.

`toPrecision(p, rnd_mode = "half-up")`

`toFixed(p, rnd_mode = "half-up")`

`toExponential(p, rnd_mode = "half-up")`

Convert the `BigDecimal` this to string with the specified precision `p`. There is no limit on the accepted precision `p`. The rounding mode can be optionally specified. `toPrecision` outputs either in decimal fixed notation or in decimal exponential notation with `p` digits of precision. `toExponential` outputs in decimal exponential notation with `p` digits after the decimal point. `toFixed` outputs in decimal notation with `p` digits after the decimal point.

6 Math mode

A new *math mode* is enabled with the "use math" directive. It propagates the same way as the *strict mode*. It is designed so that arbitrarily large integers and floating point numbers are available by default. In order to minimize the number of changes in the Javascript semantics, integers are represented either as Number or BigInt depending on their magnitude. Floating point numbers are always represented as BigFloat.

The following changes are made to the Javascript semantics:

Floating point literals (i.e. number with a decimal point or an exponent) are BigFloat by default (i.e. a `1.0` is implied). Hence `typeof 1.0 === "bigfloat"`.

Integer literals (i.e. numbers without a decimal point or an exponent) with or without the `n` suffix are BigInt if their value cannot be represented as a safe integer. A safe integer is defined as a integer whose absolute value is smaller or equal to $2^{53}-1$. Hence `typeof 1 === "number"`, `typeof 1n === "number"` but `typeof 9007199254740992 === "bigint"`.

All the bigint builtin operators and functions are modified so that their result is returned as a Number if it is a safe integer. Otherwise the result stays a BigInt.

The builtin operators are modified so that they return an exact result (which can be a BigInt) if their operands are safe integers. Operands between Number and BigInt are accepted provided the Number operand is a safe integer. The integer power with a negative exponent returns a BigFloat as result. The integer division returns a BigFloat as result.

The `^` operator is an alias to the power operator (`**`).

The power operator (both `^` and `**`) grammar is modified so that `-2^2` is allowed and yields `-4`.

The logical xor operator is still available with the `^^` operator.

The modulo operator (`%`) returns the Euclidian remainder (always positive) instead of the truncated remainder.

The integer division operator can be overloaded with `Operators.updateBigIntOperators(dictionary)`.

The integer power operator with a non zero negative exponent can be overloaded with `Operators.updateBigIntOperators(dictionary)`.